

SCHEDULING JUSTICE

A PUCAR Hackathon for FOSS United Week

A court that plans time is a national priority.

India has a pendency of 5 million cases which without reform will take 300 years for resolution. The default fix to pendency that is offered across is more judges but the evidence points elsewhere. It is not a lack of judicial capacity but a lack of how we structure, schedule and honor a judge's time that is a miss. It is a question of how we can free up a judge's time to do the judicial work.

25	63.66 lakh	~750
High Courts	cases pending across High Courts	working High Court judges today

NJDG-reported pendency as of 31 December 2025; working-judge count is our approximate estimate from sanctioned strength and vacancies.

Today, **every judge** is responsible for approximately **~8,500 cases**.

How might we enable them to **make better decisions about their time** - what to hear, when to hear it and how to schedule their day?

1. THE PROBLEM

A court day is full — but the work still does not move.

Ask anyone who has spent a day inside a court and the experience converges: litigants wait, lawyers coordinate across conflicting cases in different courtrooms without no certainty about how long a case is going to take and in what time is it going to come up for hearing, staff manage the day without a distinct plan, and judges manage a case load far larger than the hours available to hear it, while spending a large part of it on administrative aspects. It all boils down to four roles and one scheduling problem.

Litigant	“अब कोर्ट आना और जाना जीवन की कहानी बन गया है, चप्पल घिस जाती है” — going to court and back has become the story of my life; the sandals wear out from the trips.
-----------------	---

“I am told a hearing is listed today and I have to waste an entire day's work to be present. I have no idea at what point during the day the matter will be heard. So I waited the whole day. I have no visibility on what the hearing is going to be about, am I needed to do anything before I come and who needs to act for the hearing to be successful. I am dependent on my lawyer to understand every detail. Sometimes the lawyer doesn't pick up the phone and I feel quite overwhelmed and scared. The dominant feeling is hopelessness.”

Advocate	“I have 15 different matters listed in 3 different courtrooms. I keep asking my colleague what is happening in the other courtroom so just in case my matter comes up there, I am able to run and attend. Running between courtrooms seems like a key aspect of my job.”
-----------------	--

“Here, I am present in this courtroom but on my phone I have logged into the online hearing of another court to see if my case gets pulled up there.”

“I have to keep good relations with the court staff so that I can ask them to schedule my cases based on my convenience.”

“I have multiple hearings in different courts and multiple hearings at different times in the same court. Throughout the day I am scrambling for information and between courtrooms. I wish they had scheduled the hearings together so that I at least have some convenience. The matters that are old - I hardly remember the details and I show up unprepared in most hearings although I wish I didn't.”

Court Master	“Everyday I begin with a rollcall - just to see who is present and who isn't. On that day I give them some time but there is no way of knowing if it will be heard. I have to manage all the people and their requests - they have travelled all the way, I understand - I try to make it work for everyone but we are also constrained.”
---------------------	---

Judge	“I have a roster so big, every day I must be listing cases of which I am only able to hear a few and only few move forward. I am sure I can finish it but I need my time freed up and structured better. The cases are of different types and need different priorities, I need to switch context from time to time.”
--------------	---

“The lawyers are never prepared. I have cases that are so old that they need a summarization - what are the arguments made, what events have occurred, what facts have been established and what remains disputed. The lawyers can't do a thorough summarization and I waste hearings in order to just get to the same page.”

Across courts, the judge and court staff decide by hearing - how to schedule better so that cases move forward, pendency goes down and litigants and advocates are served with justice.

Scheduling is a chain of decisions:

1. Work allocation

The Chief Justice distributes cases across judges, giving each judge their own **roster** — the set of cases they are responsible for hearing.

2. Prioritization

Within their roster, a judge decides **which cases should be heard first** — based on factors such as the stage of the case, complexity, year of cognizance, urgency, and other considerations.

3. Scheduling

The judge and their court staff then decide **which cases to hear on which day**, and how many cases can realistically be packed into a day so that the maximum number of cases can be heard and moved forward. The **causelist** is the resulting schedule of hearings for the day.

For this hackathon, we will start at the **High Court, with the roster of an individual judge**. The challenge is: **given a judge's roster and a time window, how might we help them make better decisions about what to hear, when to hear it, and how to structure their day?**

2. THE POSSIBILITY

What if a court operated with a well-run schedule?

A judge's time is as precious as a surgeon's. Yet a surgeon typically knows their day in advance: how many procedures are planned, which ones need more time, what preparation needs to happen beforehand, and where their own time is most needed.

Hospitals make this possible by planning capacity across different roles, preparing in advance, and reserving scarce time for work that needs it. Other industries do this too. Airlines use historical no-show patterns to determine how many seats to sell - not underbooking or overbooking. Ride-hailing platforms continuously re-plan schedules across multiple actors as conditions change.

Courts could operate with the same discipline.

Given a judge's roster, historical data on no-shows, hearing failure rates, case stages and other relevant factors, it should be possible to generate a **better, more optimal causelist** - one that makes better use of the judge's scarce time, reduces wasted hearings, and makes the day more predictable for everyone involved.

The question is no longer just "Which cases are pending?"

It is: "Given everything we know, what is the best way to use this judge's time today?"

INSPIRATIONS FROM OTHER SECTORS

Airlines — web check-in + years of no-show data → precise overbooking, down to the seat.

Hospitals — a surgeon's day is pre-packed; prep happens before they walk in.

Restaurants — deposits and no-show fees keep a reservation book honest.

Ride-hailing — dispatch that re-packs everyone's time live as delays happen. Penalties for no-shows.

The impact

India has 25 High Courts, with sanctioned strength of over 1,100 judges. Roughly a third of those posts are vacant, leaving closer to 750 judges actually on the bench. There were 63.66 lakh cases pending across High Courts as of December 2025 which amounts to roughly 8,500 per working judge by our approximate calculation.

A scheduling tool that speeds up case disposal for one judge is therefore not just a productivity tool. Multiplied across the bench, it could become a lever for resolving millions of pending cases.

And it does not need a court-system rebuild.

District courts in Kerala, Andhra Pradesh have already shown one possible model: an independent AI transcription tool was plugged into existing court infrastructure and mandated for use. Scheduling can be another independent layer or service something courts can bolt on rather than rebuild around. This allows for easier adoption without disrupting the existing systems and actors.

3. THE CASE STUDY

3,000 pending cases over 3 months	60 hearings listed / day	20 hearings actually heard	10 effective hearings
---	------------------------------------	--------------------------------------	---------------------------------

Justice Sehgal is given a roster of 3,000 pending cases that he has to finish within 3 months of his posting in the High Court: a quarter of these cases are under a year old, while 1 in 6 are older than four years. A typical day lists 60 hearings — mostly quick mentions and interim applications, a handful of half-hour evidence sessions, and one or two hour-long final arguments.

What happens?

Out of the 30 listed hearings, 10 are actually heard. The remaining 20 are adjourned for a future date because the parties did not show up. Of the 10 that are heard, 5 are effective - meaning the purpose of hearing is met and the case moves forward towards resolution. The remaining 5 remain in the same stage because the parties were underprepared - older cases need to be summarized, facts are not clear, etc. or the events that had to happen before this hearing are still pending - warrant is pending delivery. The default gap to the next date is 60 days, whether the next step needs two minutes of work or two hours. Overall, nearly 80% of the cases are adjourned and the time of multiple actors is wasted. On a given day, the maximum number of available hours for judicial work is 7 hours. Based on the probability of a successful hearing, the causelist can be optimised to include a more realistic number of hearings.

Given a judge's roster, how might we enable them to make better decisions about what to hear, when to hear it, and how to schedule their day?

4. DESIGN FOR HOW JUDGES ACTUALLY WORK

Scheduling is not one behaviour. Across different High Courts, judges have devised their own systems and methods for listing cases. Below are three paraphrased perspectives of three HC judges who are sharing their problem statement. The goal is to design for genuine variation through configurability while keeping some design principles common and core.

Justice Sehgal — the block scheduler

When I am posted, I am given a docket of 3,000 - 4,000 cases with 2.5 months to go through them. I create predictability for the advocates and litigants by listing 10 days for every day; until the case is heard the case will continue to be listed for the same day next week. Example - I have listed 10 cases for Monday and I was only able to hear 5, the remaining 5 will be carried to the causelist next Monday until they are heard. It allows for the counsels and parties to know predictably that their case will be prioritized. It works to remove the backlog of old cases that have been pending for more than 10 years.

Second, I create predictability in my day by scheduling cases in slots based on their purpose of hearing. For example - 11am to 1:30pm is for fresh matters and notice matters. From 2:30 to 4:30, I target the oldest matters — counsel already know this. In this way I am able to manage my time well.

However, I am still only able to hear 5 cases from the 10 that I list everyday - pre-dominantly because for the older cases - the facts are not known and we waste hearing time recollecting them and ascertaining them again. This unpreparedness from the counsels on summarizing the case causes us to spend more time than required. Further, I do not get any data or insight on how I am navigating through my allocated caseload - so I have no sense of churn or indicator of my performance.

Justice Dimakar — the clusterer

My core priority is to reduce the backlog, especially across cases pending over a long time. I hear only arbitration matters and hence I classify disputes by their purpose - arguments on a certain day, appearances on a certain day. I also tend to group hearings by the advocate wherever I can. While ensuring I am admitting new matters, my key focus is every time I pass on the roster to the next judge the number of old matters have reduced significantly.

I also solve for some of the underpreparedness of the advocates in old cases by asking them to submit a front page cover before every argument hearing summarizing the key facts that have been agreed on, the critical events occurred. This helps me ensure everyone is coming to the court prepared and no one is wasting time. Re-reading an old file every twenty minutes exhausts a bench more than the number of hearings itself and hence I have mandated this submission from advocates to ensure discipline in my court.

Justice Joshi — new matters first

I'm newly posted to the bench so my focus for now is to dispose off fresh cases faster. I want to build a healthy track record so since fresh matters move faster, I take them up first. The old, complicated files need real reading time I don't have between listings, so they wait for a day I never quite get to.

This strategy ain't ideal but I tend to focus on maximising performance points by disposing of the easiest cases first. Maybe once I have gained enough seniority I would focus on the older and more complicated matters too - give them more time.

The ask

Develop a scheduling solution that helps achieve the below first principle goals. Maximising any of these goals should be at the heart of your solution. You could tackle one or many of these goals in your ideas and solutions. The meta goal is to find incentives / costs / solutions / algorithms / UX / tools that can improve efficiency and predictability of the court.

- 1. Minimise trips to court.** Litigants and counsels make several trips for their hearings. Solutions that bundle hearings into one or two longer sessions; increase asynchronous action that can be done outside a hearing; or hear different cases of the same counsel on the same day are examples of ideas that can minimize their trips to the court.
- 2. Generate an optimal causelist.** Model an algorithm that generates a causelist that caps each day to exactly the number of cases that can genuinely be heard - not overbooked or underbooked. Using the available data, optimise the causelist for hearings that if listed can be heard and if heard will move the case forward. The core parameter is efficiency - maximise for more hearings to be heard and to be substantive (case moves forward) in a given time.
- 3. Give people a real appointment.** With the causelist generated, ensure that for litigants and advocates, there is some time predictability through a defined slot — for example, 2–4pm — rather than an all-day wait. For advocates, efficiency would also improve if all their cases are listed on the same day, same slot.
- 4. Raise substantive hearing rates.** Introduce incentives / costs and solutions that increase predictability of a day in advance by gathering data on preparedness or intent to show-up from the actors. Do not list a case when its prerequisites are not met. Reduce the reasons for hearing failures. Long pending cases need quick summarization - allow counsels and parties to indicate their preparedness or generate summaries for all parties, including mandatory submissions to improve preparedness. Remind people with their next steps or pending tasks.
- 5. Make the next date meaningful.** Similarly, recommend the most optimal next hearing date at the end of a hearing. Match the gap to the actual procedural minimum, rather than a flat default. Based on what happened in today's hearing and the next purpose of hearing, predict the right amount of time required for the case to move forward.
- 6. Make scheduling configurable for judges.** Allow judges to configure their rules to provide flexibility. Support blocks, clustering and ageing-case quotas without changing code. However, keep some goals unchangeable - example: ageing cases getting deprioritised shouldn't be configurable.
- 7. Data is king.** Generate insights that allow judges to make better decisions about their scheduling. Allow them to have visualisations that help them understand the costs / effect of their rules overriding the most optimal scheduling. Flag ageing cases, repeat adjournments, and cases stuck at one stage too long. Allow them to simulate best ways of scheduling for maximum efficiency and predictability.

5. WHAT YOU CAN BUILD WITH

What we give you

- Judge's [roster sample of 100 cases](#); do use a script to generate more cases for the 3000 roster
- [Court calendar](#) for holidays, separately account for the judge's personal leaves
- [Sample causelist](#) (present - you can create better causelists)
- A [hearing-type reference](#) table with minimum-maximum hearings in a case, time to complete in a day and ideal gap to next date.
- A [hearing type substantive-ness probability](#) showing what is the likelihood of substantive hearings for a given purpose of hearing.
- A [hearing-failure reasons](#) table showing why listed hearings do not happen or move forward.

What does this data contain - how do I understand this?

You are not starting with a blank slate. The datasets give you a small but realistic picture of a court's docket: **100 cases, where each case is in its lifecycle, what has happened to it so far, and what typically happens at each kind of hearing.** The data is not here simply to reproduce what happens today. It gives you the **starting conditions and observed patterns** from which you can build a scheduling model and then test what happens when you change the way the court operates.

1. The case lifecycle: 11 stages

Every case in this docket moves through the same fixed sequence of substantive stages, from filing to verdict. A case can sit at a stage for multiple hearings (adjournments, notice returns, absent parties) before it moves to the next one.

#	Stage	What happens here
1	Admission	The court examines the complaint (affidavit under s.225 BNSS), checks there's no undue delay or resolves jurisdictional questions, and decides if there is sufficient ground to take the case on file.
2	Delay Condonation Hearing	If the complaint was filed late, the court hears the accused's objections and decides whether to excuse ("condone") the delay before the case can proceed at all.
3	Cognizance	The court formally takes cognizance of the offence — it is satisfied there's sufficient ground to proceed against the accused, and the case is registered.
4	Appearance	Summons are issued and the court tracks whether the accused has appeared. Hearings here are often about summons being served, returned unclaimed, or needing to be reissued.
5	Warrant	If the accused still doesn't appear despite summons, the court issues a warrant (bailable or non-bailable) to compel appearance.
6	Plea	The accused appears, the particulars of the offence are read out and explained, and the accused pleads guilty or not guilty. A not-guilty plea sends the case to trial.
7	Examination Under S351 BNSS	The court examines the accused directly to give them a chance to explain the circumstances against them, and the accused can file a list of defence witnesses.
8	Evidence Complainant	The complainant's side presents its evidence — examination-in-chief and cross-examination.
9	Evidence Accused	The defence presents its evidence in the same way.
10	Arguments	Once evidence closes, both sides make their final oral arguments.
11	Judgement	The court pronounces its verdict — conviction or acquittal — and, on conviction, sentence and/or compensation. Besides judgment, the parties can also withdraw the case if they reach a settlement through negotiations outside the court or through alternate dispute resolution mechanisms.

Besides these, three hearing types sit outside the fixed sequence which means they can interrupt a case at almost any point rather than occupying a fixed slot in it:

Type	What it is	When it can occur
Bail	The accused's application to be released on bail.	Any point before trial starts — once the case reaches Evidence Complainant or later, the accused's bail status is already settled, so it can't recur.
Reports	Progress updates — most commonly mediation referrals reporting back to court.	Any stage.
Application Review	Miscellaneous/interlocutory applications (CMPs) filed by either party — objections, extensions, procedural requests.	Any stage.

2. What each dataset tells you:

A. Start with the 100 cases

The **Sample 100 Cases** is your starting point. For each case, you can see things like:

- when the case was filed — which tells you its **age**
- what stage of the case it is currently at
- what happened at its last hearing
- who was present
- what the next hearing is for
- how many hearings of each type the case has already gone through

Think of this as the **state of the docket** at a particular point in time. The 100 cases are only a sample. Your scheduler should be able to work with a much larger roster of 3000 cases given to Justice Sehgal — you can use a script to generate more cases from the patterns in the data.

One important simplification for the hackathon: you can treat the roster as your starting point. Don't worry about preserving the existing next hearing dates. Imagine the judge starts with this roster and has to decide **what to hear, when to hear it and how to use the available time which is 2.5 months**.

B. A hearing type tells you something about the work involved

A case does not simply have "a hearing". It moves through different stages or purposes — Admission, Appearance, Warrant, Plea, Evidence, Arguments, etc. — and each type behaves differently. The **Hearing Type Reference** gives you the observed range and typical number of hearings a case has at each stage, based on the pilot court's real case history. For example, Evidence Complainant may require several hearings, while Admission may usually be completed in fewer. The table also gives you information such as the expected time required for different hearing types and the appropriate gap before the next hearing. This gives you a sense of the **shape and cost of different kinds of work** on a judge's calendar.

But these are not fixed values that every case must follow. **A hearing duration is not guaranteed**. You could, for example, sample duration from a distribution based on the observed data, so that the time taken by a hearing can vary from case to case.

C. Not every listed hearing actually happens and not every hearing that happens moves the case forward

This is an important part of the data. A hearing is called substantive if it moves the case forward - meaning it goes into the next purpose of hearing. The **Nonsubstantive Reasons by Hearing Type** table shows why a hearing may fail to happen or fail to move the case forward. Reasons include things like:

- a party not appearing
- an adjournment being sought
- a court holiday
- waiting for a summons or warrant to be returned
- other prerequisites not being completed

These probabilities are estimates grounded in the real court data. So, for example, if the data suggests that a certain percentage of people do not appear, that does not necessarily mean **every simulation should simply apply the same percentage mechanically**. You could model attendance as a probability or introduce behavioural assumptions that allow different actors or cases to behave differently.

The **Sample Causelist** gives you another view of this. It is an actual day's cause list from the pilot court, including the attendance and order text recorded that day. Use it to understand what a real court day actually looks like — rather than assuming that everything on a cause list gets heard exactly as planned.

D. The calendar tells you when the court can actually work

The **Court Calendar** gives you the working days, weekly offs and gazetted holidays for the simulation period. This is the basic constraint around which everything else has to fit. You should separately account for the judge's personal leave where relevant.

The question is then not just "what cases should be scheduled?" but "what can realistically fit into the time available?"

E. Put the datasets together

The interesting part is not any one table. It is what happens when you connect them. For a given case, you can start asking:

How old is it? → Where is it in its lifecycle? → What kind of hearing does it need? → How long might that hearing take? → How likely is it to actually happen? → If it happens, how likely is it to move the case forward? → When should it come back?

Across the roster, you can then ask:

Given 3000 cases, the available 3 months court time and the way hearings actually behave, what should this judge's schedule for the next few weeks look like?

There are other signals you can choose to model too: advocate and party identifiers can help you identify cases that might sensibly be clustered; repeated adjournments can reveal cases that are getting stuck; and case age can help you see what happens to older cases as different scheduling rules are applied.

F. The data is deliberately incomplete: This is where you can model

The datasets do not tell you exactly what the answer is. Some fields are observed from real court data. Others are estimates or generated from those observations. Some things that would matter in a real court — such as real-time availability, preparedness or a judge's personal scheduling preferences — are not fully captured.

You can use the data provided, generate additional cases or variables from the observed distributions, and introduce your own assumptions where you think they are useful. **Make those assumptions visible**. More importantly, think about **what you are actually changing in your model and whether that change produces a material difference**.

For example:

- What happens if you change how cases are prioritised?
- What happens if you change how much of the judge's available time you fill?
- What happens if you cluster cases by advocate?
- What happens if you change the gap to the next hearing based on the actual work required, rather than using a standard gap?
- What happens if hearing duration is sampled rather than treated as fixed?
- What happens if parties or advocates have the ability to behave differently — for example, to appear, seek an adjournment or respond to a change in schedule?
- What happens when the schedule changes because a hearing takes longer than expected or a party does not show?

These are not necessarily features you must build. They are examples of **variables, policies or behaviours that can be shifted and tested**. A useful submission should be able to show not just *what schedule it produces*, but **what changed in the model, why you changed it, and what happen**

To begin with, the tool could:

- Take in a judge's roster and the distribution for hearing failures.
- Design an optimal scheduling algorithm to pack the roster in a limited time frame

- Let the judge set their own rules — priority, time budget, and other preferences.
- Generate insights and visualisations for judges to make data-backed decisions.
- Generate a daily cause list based on final set rules.
- Capture new data about the characteristics of a case, the parties, their advocates for better behavioral modelling.
- Suggest the next-best date when a hearing does not conclude.

New data you can collect / generate - examples:

- Next hearing purpose — mention, evidence, arguments, admission.
- Accurate hearing duration — from the hearing-type reference table.
- Real-time updating adjournment / no-show likelihood.

6. HOW WILL WE KNOW IT WORKS?

A complete solution is not just an algorithm. **A judge has to be able to use it, understand it, trust it, and shape it to the way their court works.** We will therefore look at two things:

1 · Scheduling Performance

The scheduling algorithm is at the heart of this hackathon. The more your solution can account for **dynamic, real-world and behavioural factors**, the more sophisticated the scheduling model becomes. You should be able to explain how you arrived at your approach using the data provided and walk us through the key decisions built into your algorithm. Ultimately, the schedule should optimise for:

What we measure	What this means
Utilisation	Is the judge's day well packed — without being overbooked or leaving significant capacity unused?
Predictability	Is the schedule realistic? Are cases actually heard when they are listed and is enough time available based on their purpose?
Substantiveness	Is court time being used effectively? If a case is heard, does it meaningfully advance towards judgment or resolution?
Backlog-age impact	When the schedule runs forward, does it reduce the number of cases sitting in the 5+ year buckets? Cases with high ages should not be deprioritised.
Next-date recommendation	Given everything known about a case, does the system recommend a sensible and efficient next date?

2 · Visualise Insight

A good scheduling system should not simply **produce a causelist and execute it silently**. It should help a judge and their staff **see, understand and shape the consequences of their scheduling decisions**. Different judges work differently and may choose to override what the algorithm recommends. That should be possible but the system should help them understand the trade-offs.

For example:

- What happens to utilisation if I move these 10 cases?
- What happens to predictability if I overbook this day?
- Which cases are at risk of ageing further?
- Which cases have had repeated adjournments?
- Where is my docket beginning to drift?
- What happens to my 3+ and 4+ year backlog if I make this change?
- What will the next few weeks of my court look like if I choose this option?

The goal is to give judges a **live picture of their court**, rather than simply handing them a schedule.

Think of it like running a tight hospital: you can see capacity, bottlenecks, ageing cases and risks early, make informed decisions, and understand the consequences of changing the plan. The strongest solutions will therefore not just **schedule the court** — they will help a judge **understand and run the court**.

Submission Complexity

The scoring will consider **how much real-world complexity your solution is able to capture**. Think of the levels as a progression:

L1 > Easy: Fixed behaviour

You build the basic scheduling logic and show how it works.

For example:

- A simple scheduling or optimisation algorithm
- Fixed rules for how cases are prioritised and scheduled
- A walkthrough of the algorithm
- A simple visualisation of the resulting schedule

In short: You are modelling the **schedule**.

L2 > Medium: Dynamic + realistic data

You introduce changing conditions, realistic data and **costs or incentives** into the scheduling model.

For example:

- Model the **distribution of different variables** — such as hearing duration, adjournment frequency, case type, lawyer availability or time between hearings — rather than assuming every case behaves the same way
- Use these distributions to make the model and scheduling decisions more accurate
- Allow the schedule to change based on new information or decisions
- Introduce different **costs, constraints or incentives** and model their impact on scheduling decisions
- Allow a judge to override a scheduling rule and see the impact on efficiency and predictability
- Use static or dynamic visualisations to show how different decisions, costs or incentives affect the court
- Integrate a small classification or prediction model into the scheduling logic

In short: You are modelling a **changing and more realistic court**, including the trade-offs that shape decisions.

L3 > Hard: Behavioural / Agent-based

You introduce **people who behave** into the model. Instead of assuming that everyone follows fixed rules or statistical distributions, you model actors as agents who can make decisions.

For example:

- Create an agent representing a litigant, lawyer or judge
- Give the agent simple characteristics or a “personality”
- Use a simple agentic/behavioural model — for example, a JEV-style tool or small model that can be asked to make a decision
- The agent could decide whether they will appear for a hearing, seek an adjournment, or respond to a scheduling change
- Model how agents respond to different costs, incentives and constraints
- Simulate how those individual decisions affect the court's schedule, and how the changing schedule affects future decisions

The agent does **not** need to be sophisticated. Even a simple behavioural agent that makes a decision based on the situation can demonstrate this level of complexity.

In short: You are modelling the **people inside the court**.

The ambition

The progression we are looking for is: **L1 — Model the schedule → L2 — Model realistic variation and trade-offs → L3 — Model the people who respond to them.**

The longer-term possibility is much bigger: **a simulation of a court where judges, lawyers and litigants each behave as agents, with different characteristics, incentives and constraints allowing us to simulate how a real court might behave before we change it.**

For this hackathon, you don't need to build the whole simulation. **Show us how far you can take the idea but when in doubt - remember to simply simplify!**

7. BUILD SOMETHING WE CAN TAKE TO COURT

Start here: README + GitHub + .md files.

Use the time and space to code, ask questions, test ideas and dig deeper into the problem. By 5 PM, your code must be uploaded to a branch and ready for merge. You'll then have 10 minutes to present your prototype to the panel. The panel will use the criteria above to identify the two strongest teams / individuals. Their code will be taken forward for integration into the DRISTI 2.0 stack as the results are declared.

DRISTI 2.0 is going live across 3 states in India.

And more: The selected teams / individuals will work directly with the Court Master and Judge to test, tweak and finish the tool in a live court context. The final tool will be released as open source and, if useful, taken to other High Courts.

So don't just build something that works for the hackathon. Build something that could change how a court works.